

# Προχωρημένα Θέματα Βάσεων Δεδομένων

Μαρία Κοιλαλού | Μυρτώ Ορφανάκου  
AM:03119211 | AM:03119173

15 Ιανουαρίου 2024

1

Για την εγκατάσταση του Apache Spark ώστε να εκτελείται πάνω από το διαχειριστή πόρων του Apache Hadoop και YARN, δημιουργήσαμε δύο Virtual Machines στην πλατφόρμα okeanos-knossos ακολουθώντας τις οδηγίες που μας δόθηκαν. Οι web εφαρμογές είναι προσβάσιμες και διαμορφωμένες σύμφωνα με τις οδηγίες στα παρακάτω links:

YARN  
HDFS  
Spark History

## Spark Session

Αρχικά, υλοποιήθηκε μια συνάρτηση που δημιουργεί ένα Spark Session και δέχεται ως όρισμα τον αριθμό των executors.

```
1 from pyspark.sql import SparkSession
2
3 def create_spark_session(num_executors):
4
5     spark = SparkSession.builder \
6         .appName("ADV DATABASES") \
7         .config("spark.executor.instances", str(num_executors)) \
8         .config("spark.executor.cores", "1") \
9         .config("spark.executor.memory", "1g") \
10        .config("spark.driver.memory", "4g") \
11        .config("spark.cleaner.periodicGC.interval", "1min") \
12        .config("spark.memory.offHeap.enabled", "true") \
13        .config("spark.memory.offHeap.size", "1g") \
14        .config("spark.default.parallelism", "4") \
15        .config("spark.serializer", "org.apache.spark.serializer.KryoSerializer") \
16        .config("spark.dynamicAllocation.enabled", "true") \
17        .config("spark.dynamicAllocation.minExecutors", "2") \
18        .config("spark.dynamicAllocation.maxExecutors", "4") \
19        .getOrCreate()
20
21     return spark
```

Κάθε φορά που θέλουμε να ξεκινήσουμε νέο session καλούμε την συνάρτηση `create_spark_session()` με τον αντίστοιχο αριθμό executors που θέλουμε.

```
1 spark = create_spark_session(4)
2 print("spark session created")
```

## 2 Create DataFrame

Δημιουργήσαμε ένα schema για το DataFrame του βασικού συνόλου δεδομένων. Τα αρχικά ονόματα των στηλών διατηρήθηκαν, ενώ οι τύποι δεδομένων προσαρμόστηκαν με βάση τα ζητούμενα.

Στη συνέχεια δημιουργήθηκε ένα ενιαίο DataFrame για όλα τα δεδομένα, όπως φαίνεται στον παρακάτω κώδικα:

```
1 schema1 = "`DR_NO` STRING, \  
2 `Date Rptd` STRING, \  
3 `DATE OCC` STRING, \  
4 `TIME OCC` INTEGER, \  
5 `AREA` INTEGER, \  
6 `AREA NAME` STRING, \  
7 `Rpt Dist No` INTEGER, \  
8 `Part 1-2` INTEGER, \  
9 `Crm Cd` INTEGER, \  
10 `Crm Cd Desc` STRING, \  
11 `Mocodes` STRING, \  
12 `Vict Age` INTEGER, \  
13 `Vict Sex` STRING, \  
14 `Vict Descent` STRING, \  
15 `Premis Cd` INTEGER, \  
16 `Premis Desc` STRING, \  
17 `Weapon Used Cd` INTEGER, \  
18 `Weapon Desc` STRING, \  
19 `Status` STRING, \  
20 `Status Desc` STRING, \  
21 `Crm Cd 1` INTEGER, \  
22 `Crm Cd 2` INTEGER, \  
23 `Crm Cd 3` INTEGER, \  
24 `Crm Cd 4` INTEGER, \  
25 `LOCATION` STRING, \  
26 `Cross Street` STRING, \  
27 `LAT` DOUBLE, \  
28 `LON` DOUBLE"  
29  
30  
31 data1 = spark.read.csv("/user/ubuntu/ta/advanced-db/data/crime_data_2010.csv", header=True, schema=schema1)  
32 data2 = spark.read.csv("/user/ubuntu/ta/advanced-db/data/crime_data_2020.csv", header=True, schema=schema1)  
33  
34 df = data1.union(data2).distinct()  
35  
36 df = df.withColumn("Date Rptd", to_date(col("Date Rptd"), "MM/dd/yyyy hh:mm:ss a")) \  
37 .withColumn("DATE OCC", to_date(col("DATE OCC"), "MM/dd/yyyy hh:mm:ss a"))  
38  
39 df.count()  
40 print(f"Total number of rows: {df.count()}")  
41  
42 df.printSchema()
```

Ο συνολικός αριθμός γραμμών του συνόλου δεδομένων, καθώς και ο τύπος κάθε στήλης είναι τα εξής:

```
Total number of rows: 2913595  
root  
|-- DR_NO: string (nullable = true)  
|-- Date Rptd: date (nullable = true)  
|-- DATE OCC: date (nullable = true)  
|-- TIME OCC: integer (nullable = true)  
|-- AREA: integer (nullable = true)  
|-- AREA NAME: string (nullable = true)  
|-- Rpt Dist No: integer (nullable = true)  
|-- Part 1-2: integer (nullable = true)  
|-- Crm Cd: integer (nullable = true)  
|-- Crm Cd Desc: string (nullable = true)  
|-- Mocodes: string (nullable = true)  
|-- Vict Age: integer (nullable = true)  
|-- Vict Sex: string (nullable = true)  
|-- Vict Descent: string (nullable = true)  
|-- Premis Cd: integer (nullable = true)
```

```

|-- Premis Desc: string (nullable = true)
|-- Weapon Used Cd: integer (nullable = true)
|-- Weapon Desc: string (nullable = true)
|-- Status: string (nullable = true)
|-- Status Desc: string (nullable = true)
|-- Crm Cd 1: integer (nullable = true)
|-- Crm Cd 2: integer (nullable = true)
|-- Crm Cd 3: integer (nullable = true)
|-- Crm Cd 4: integer (nullable = true)
|-- LOCATION: string (nullable = true)
|-- Cross Street: string (nullable = true)
|-- LAT: double (nullable = true)
|-- LON: double (nullable = true)

```

### 3 Query 1

Υλοποιήθηκε μία συνάρτηση για το Query 1 χρησιμοποιώντας DataFrame API:

#### Dataframe API

```

1 def query1_df(df):
2     crime_date = df.withColumn("Year", year("DATE OCC")).withColumn("Month", month("DATE OCC"))
3
4     count = crime_date.groupBy("Year", "Month").count()
5
6     window_spec = Window.partitionBy("Year").orderBy(desc("count"))
7     top_months = count.withColumn("rank", dense_rank().over(window_spec)).filter(col("rank") <= 3)
8
9     top_months = top_months.orderBy("Year", "rank")
10
11     return top_months

```

Στη συνέχεια το Query 1 υλοποιήθηκε με SQL API:

#### SQL API

```

1 def query1_sql(df):
2     crime_date = df.withColumn("Year", year("DATE OCC")).withColumn("Month", month("DATE OCC"))
3
4     # Δημιουργία προσωρινής προβολής
5     crime_date.createOrReplaceTempView("crimes")
6
7     # SQL ερώτημα για την εύρεση των τριών μηνών με τον υψηλότερο αριθμό εγκλημάτων ανά έτος
8     query1 = """
9     SELECT Year, Month, count, rank
10    FROM (
11        SELECT Year, Month, count(*) AS count,
12               DENSE_RANK() OVER (PARTITION BY Year ORDER BY count(*) DESC) AS rank
13        FROM crimes
14        GROUP BY Year, Month
15    )
16    WHERE rank <= 3
17    ORDER BY Year, rank
18    """
19
20     top_months = crime_date.sparkSession.sql(query1)
21
22     return top_months

```

Παρακάτω φαίνονται τα αποτελέσματα του Query 1 για τις δύο διαφορετικές υλοποιήσεις, καθώς και οι αντίστοιχοι χρόνοι εκτέλεσης. Παρατηρούμε ότι υπάρχει διαφορά στην επίδοση των δύο APIs. Συγκεκριμένα, το DataFrame API πραγματοποιήθηκε σε πολύ λιγότερο χρόνο σε σχέση με το SQL API.

```

+----+-----+-----+----+
|Year|Month|count|rank|
+----+-----+-----+----+
|2010| 1|19515| 1|
|2010| 3|18131| 2|
|2010| 7|17856| 3|
|2011| 1|18134| 1|
|2011| 7|17283| 2|
|2011|10|17034| 3|
|2012| 1|17943| 1|
|2012| 8|17661| 2|
|2012| 5|17502| 3|
|2013| 8|17440| 1|
|2013| 1|16820| 2|
|2013| 7|16644| 3|
|2014| 7|12196| 1|
|2014|10|12133| 2|
|2014| 8|12028| 3|
|2015|10|19219| 1|
|2015| 8|19011| 2|
|2015| 7|18709| 3|
|2016|10|19659| 1|
|2016| 8|19490| 2|

```

only showing top 20 rows

Q1 Dataframe time: 0.3153994083404541 seconds.

```

+----+-----+-----+----+
|Year|Month|count|rank|
+----+-----+-----+----+
|2010| 1|19515| 1|
|2010| 3|18131| 2|
|2010| 7|17856| 3|
|2011| 1|18134| 1|
|2011| 7|17283| 2|
|2011|10|17034| 3|
|2012| 1|17943| 1|
|2012| 8|17661| 2|
|2012| 5|17502| 3|
|2013| 8|17440| 1|
|2013| 1|16820| 2|
|2013| 7|16644| 3|
|2014| 7|12196| 1|
|2014|10|12133| 2|
|2014| 8|12028| 3|
|2015|10|19219| 1|
|2015| 8|19011| 2|
|2015| 7|18709| 3|
|2016|10|19659| 1|
|2016| 8|19490| 2|

```

only showing top 20 rows

Q1 SQL time: 0.6481153964996338 seconds.

## 4 Query 2

### DataFrame\SQL API

Υλοποιήθηκε μία συνάρτηση για το Query 2 χρησιμοποιώντας DataFrame/ SQL API:

```
1 def query2_df(df):
2
3
4 def day_part(hour):
5     if 500 <= hour < 1200:
6         return "Πρωί"
7     elif 1200 <= hour < 1700:
8         return "Απόγευμα"
9     elif 1700 <= hour < 2100:
10        return "Βράδυ"
11    else:
12        return "Νύχτα"
13
14 day_part_udf = udf(day_part, StringType())
15
16 df_day_part = df.withColumn("DayPart", day_part_udf(col("TIME OCC")))
17
18 df_street_crimes = df_day_part.filter(col("Premis Desc") == "STREET").groupBy("DayPart").count().orderBy(col("count").desc())
19
20 return df_street_crimes
```

### RDD API

Στη συνέχεια το Query 2 υλοποιήθηκε με RDD API:

```
1 def query2_rdd(df):
2
3
4 def day_part(hour):
5     if 500 <= hour < 1200:
6         return "Πρωί"
7     elif 1200 <= hour < 1700:
8         return "Απόγευμα"
9     elif 1700 <= hour < 2100:
10        return "Βράδυ"
11    else:
12        return "Νύχτα"
13
14 rdd = df.rdd.filter(lambda row: row['Premis Desc'] == 'STREET')
15
16 def map_day_part(record):
17     hour = int(record["TIME OCC"])
18     part = day_part(hour)
19     return (part, 1)
20
21 rdd_mapped = rdd.map(map_day_part)
22 rdd_reduced = rdd_mapped.reduceByKey(lambda a, b: a + b)
23
24 rdd_street_crimes = rdd_reduced.sortBy(lambda x: x[1], ascending=False)
25
26 return rdd_street_crimes
```

Ακολουθούν τα αποτελέσματα του Query 2 για τις δύο διαφορετικές υλοποιήσεις, καθώς και οι αντίστοιχοι χρόνοι εκτέλεσης. Παρατηρούμε ότι υπάρχει διαφορά στην επίδοση των δύο APIs. Συγκεκριμένα, το DataFrame/ SQL API πραγματοποιήθηκε σε πολύ λιγότερο χρόνο σε σχέση με το RDD API.

```
+-----+-----+
| DayPart| count|
+-----+-----+
| Νύχτα|231546|
| Βράδυ|182141|
| Απόγευμα|143974|
```

```
| Πρωί120358|
```

```
+-----+-----+
```

Q2 Dataframe time: 0.31096601486206055 seconds.

[(Νύχτα, 231546), (Βράδυ, 182141), (Απόγευμα, 143974), (Πρωί, 120358)]

Q2 RDD time: 39.22880935668945 seconds.

## 5 Query 3

Στο συγκεκριμένο ερώτημα χρησιμοποιήθηκαν τα δευτερεύοντα σύνολα δεδομένων `revgecoding` και `LA_income_2015`. Επιπλέον, πραγματοποιήθηκε `map` για τα `Vict Descent`.

```
1 data3 = spark.read.csv("/user/ubuntu/ta/advanced-db/data/LA_income_2015.csv", header=True, schema=schema2)
2 data4 = spark.read.csv("/user/ubuntu/ta/advanced-db/data/revgecoding.csv", header=True, schema=schema3)
3
4 schema3 = "`LAT` DOUBLE, \
5           `LON` DOUBLE, \
6           `ZIPcode` INTEGER"
7
8 schema2 = "`Zip Code` INTEGER, \
9           `Community` STRING, \
10          `Estimated Median Income` STRING"
11
12 descent_mapping = {
13     'A': 'Other Asian',
14     'B': 'Black',
15     'C': 'Chinese',
16     'D': 'Cambodian',
17     'F': 'Filipino',
18     'G': 'Guamanian',
19     'H': 'Hispanic/Latin/Mexican',
20     'I': 'American Indian/Alaskan Native',
21     'J': 'Japanese',
22     'K': 'Korean',
23     'L': 'Laotian',
24     'O': 'Other',
25     'P': 'Pacific Islander',
26     'S': 'Samoan',
27     'U': 'Hawaiian',
28     'V': 'Vietnamese',
29     'W': 'White',
30     'X': 'Unknown',
31     'Z': 'Asian Indian'
32 }
33
34 data3 = data3.withColumn("Estimated Median Income", regexp_replace(col("Estimated Median Income"), "\\$", ""))
35 data3 = data3.withColumn("Estimated Median Income", regexp_replace(col("Estimated Median Income"), ",", "").cast("float"))
36
37 crime_year = df.withColumn("Year", year("DATE OCC"))
38
39 crime_2015 = crime_year.filter(
40     (col("Year") == 2015) &
41     (col("Vict Descent").isNotNull()))
42
43 def map_descent(code):
44     return descent_mapping.get(code, "Unknown") # Default to "Unknown" if code not found
45
46 map_descent_udf = udf(map_descent, StringType())
47
48 crime_2015 = crime_2015.withColumn("Vict Descent", map_descent_udf(crime_2015["Vict Descent"]))
49
50 revgecoding = data4.dropDuplicates(['LAT', 'LON'])
```

Έπειτα υλοποιήθηκε το Query 3 χρησιμοποιώντας DataFrame/ SQL API.

```
1 def query3 (crime_2015, data3, revgecoding):
```

```

2 crime_zip = crime_2015.join(revgecoding, ["LAT", "LON"], "left")
3
4 best3_zip = data3.orderBy("Estimated Median Income", ascending=False).limit(3)
5 worst3_zip = data3.orderBy("Estimated Median Income", ascending=True).limit(3)
6
7 best3_zip_list = [row['Zip Code'] for row in best3_zip.collect()]
8 worst3_zip_list = [row['Zip Code'] for row in worst3_zip.collect()]
9
10 crimes = crime_zip.filter(
11     (col("ZIPcode").isin(best3_zip_list)) |
12     (col("ZIPcode").isin(worst3_zip_list))
13 )
14
15 vict_descent_count = crimes.groupBy("Vict Descent").count().orderBy("count", ascending=False)
16
17
18
19 return vict_descent_count

```

To Query 3 εκτελέστηκε σε ένα for loop, δημιουργώντας διαδοχικά τρία Spark Sessions για 2, 3 και 4 executors. Παρακάτω φαίνονται τα αποτελέσματα και η διάρκεια της κάθε εκτέλεσης αντίστοιχα. Παρατηρούμε ότι ταχύτερα εκτελέστηκε το Query με 3 executors, έπειτα με 2 και τέλος με 4.

## 2 Executors

```

+-----+-----+
|      Vict Descent|count|
+-----+-----+
|Hispanic/Latin/Me...| 1053|
|          White| 610|
|          Black| 349|
|          Other| 272|
|          Unknown| 71|
|      Other Asian| 46|
|          Korean| 4|
|          Chinese| 1|
|American Indian/A...| 1|
+-----+-----+

```

Number of Executors: 2  
Q3 time: 10.069442987442017 seconds.

## 3 Executors

```

+-----+-----+
|      Vict Descent|count|
+-----+-----+
|Hispanic/Latin/Me...| 1053|
|          White| 610|
|          Black| 349|
|          Other| 272|
|          Unknown| 71|
|      Other Asian| 46|
|          Korean| 4|
|          Chinese| 1|
|American Indian/A...| 1|
+-----+-----+

```

Number of Executors: 3  
Q3 time: 5.2063148021698 seconds.

## 4 Executors

```
+-----+-----+
|      Vict Descent|count|
+-----+-----+
|Hispanic/Latin/Me...| 1053|
|      White|    610|
|      Black|   349|
|      Other|   272|
|      Unknown|   71|
|      Other Asian|   46|
|      Korean|    4|
|      Chinese|    1|
|American Indian/A...|    1|
+-----+-----+
```

Number of Executors: 4

Q3 time: 12.506168365478516 seconds.

## 6 Query 4

Για το ζητούμενο 6 χρησιμοποιήθηκαν επιπλέον τα δεδομένα [LAPD\\_Police\\_Stations](#). Το Query 4 υλοποιήθηκε με DataFrame/SQL API και εκτελέστηκε με 4 executors.

Αρχικά δημιουργήθηκε μια συνάρτηση για τον υπολογισμό της απόστασης μεταξύ 2 σημείων με συγκεκριμένες συντεταγμένες. Στη συνέχεια πραγματοποιήθηκαν δύο join, ένα για κάθε ερώτημα

Έπειτα, για την περίπτωση του κοντινότερου station, πραγματοποιήθηκε ένα cross join μεταξύ των DataFrame των εγκλημάτων και του [LAPD\\_Police\\_Stations](#) για τον εντοπισμό του κοντινότερου τμήματος σε κάθε έγκλημα.

```
1 schema4 = "`X` DOUBLE, \
2           `Y` DOUBLE, \
3           `FID` INTEGER, \
4           `DIVISION` STRING, \
5           `LOCATION` STRING, \
6           `PREC` INTEGER"
7
8 data5 = spark.read.csv("/user/ubuntu/ta/advanced-db/data/LAPD_Police_Stations.csv", header=True, schema=schema4)
```

```
1 def query4(df, data5):
2
3     def haversine(lat1, lon1, lat2, lon2):
4         # Radius of the Earth in kilometers
5         R = 6371.0
6
7         lat1_rad = math.radians(lat1)
8         lon1_rad = math.radians(lon1)
9         lat2_rad = math.radians(lat2)
10        lon2_rad = math.radians(lon2)
11
12        dlat = lat2_rad - lat1_rad
13        dlon = lon2_rad - lon1_rad
14
15        a = math.sin(dlat / 2)**2 + math.cos(lat1_rad) * math.cos(lat2_rad) * math.sin(dlon / 2)**2
16        c = 2 * math.atan2(math.sqrt(a), math.sqrt(1 - a))
17
18        distance = R * c
19        return distance
20
21    def get_distance(lat1, long1, lat2, long2):
22
23        def is_valid_coordinate(lat, lon):
24            return -90 <= lat <= 90 and -180 <= lon <= 180
25
26
27        if not is_valid_coordinate(lat1, long1) or not is_valid_coordinate(lat2, long2):
28            # Print the invalid rows
```



```

29         print(f"Invalid row: lat1={lat1}, long1={long1}, lat2={lat2}, long2={long2}")
30         return -1
31
32     try:
33         return haversine(lat1, long1, lat2, long2)
34     except ValueError:
35         return -1
36
37 df = df.filter(
38     (df["AREA NAME"] != "Null Island") &
39     (df["Weapon Used Cd"].substr(1, 1) == "1")
40 )
41
42 joined_df = df.join(data5, df["AREA"] == data5["PREC"])
43
44 distance_udf = udf(get_distance)
45
46 distance_df = joined_df.withColumn(
47     "DISTANCE",
48     distance_udf(
49         F.col("LAT"), F.col("LON"),
50         F.col("Y"), F.col("X")
51     ).cast("double")
52 )
53
54 query_4_1a = distance_df.groupBy("Year").agg(
55     F.avg("DISTANCE").alias("average_distance"),
56     F.count("*").alias("#")
57 ).orderBy("Year").withColumnRenamed("Year", "year")
58
59 query_4_1b = distance_df.groupBy("DIVISION").agg(
60     F.avg("DISTANCE").alias("average_distance"),
61     F.count("*").alias("#")
62 ).orderBy(F.desc("#")).withColumnRenamed("DIVISION", "division")
63
64 printΑπόσταση(" από το αστυνομικό τμήμα που ανέλαβε την έρευνα για το περιστατικό:")
65 print("(a)")
66 query_4_1a.show()
67 print("(b)")
68 query_4_1b.show()
69
70 cross_joined_df = df.crossJoin(data5.withColumnRenamed("LAT", "Y").withColumnRenamed("LON", "X"))
71
72 cross_joined_df = cross_joined_df.withColumn(
73     "DISTANCE",
74     distance_udf(col("LAT"), col("LON"), col("Y"), col("X")).cast("double")
75 )
76
77 windowSpec = Window.partitionBy("DR_NO").orderBy("DISTANCE")
78
79 nearest_station_df = cross_joined_df.withColumn(
80     "row_num",
81     F.row_number().over(windowSpec)
82 ).filter(col("row_num") == 1).drop("row_num")
83
84 query_4_2a = nearest_station_df.groupBy("Year").agg(
85     F.avg("DISTANCE").alias("average_distance"),
86     F.count("*").alias("#")
87 ).orderBy("Year").withColumnRenamed("Year", "year")
88
89 query_4_2b = nearest_station_df.groupBy("DIVISION").agg(
90     F.avg("DISTANCE").alias("average_distance"),
91     F.count("*").alias("#")
92 ).orderBy(F.desc("#")).withColumnRenamed("DIVISION", "division")
93
94
95 printΑπόσταση(" από το πλησιέστερο αστυνομικό τμήμα:")
96 print("(a)")
97 query_4_2a.show()
98 print("(b)")
99 query_4_2b.show()

```

Απόσταση από το αστυνομικό τμήμα που ανέλαβε την έρευνα για το περιστατικό:

(a)

+-----+-----+-----+		
+-----+-----+-----+		
Year	num_crimes	average_distance
+-----+-----+-----+		
2010	8213	4.315547525861609
2011	7232	2.7931783031826134
2012	6550	37.401521647671025
2013	5838	2.826412721201962
2014	4230	11.631025289489838
2015	6763	2.70609799276239
2016	8100	2.7176445421299724
2017	7788	5.955847913803834
2018	7413	2.732823649229879
2019	7129	2.7399419721721476
2020	8491	8.614767812336167
2021	9767	30.97834129556094
2022	10025	2.60865618645079
2023	8741	2.5551410574543145
+-----+-----+-----+		

(b)

+-----+-----+-----+		
+-----+-----+-----+		
DIVISION	num_crimes	average_distance
+-----+-----+-----+		
77TH STREET	94474	13.162079052889169
SOUTHEAST	72832	14.527525557922345
SOUTHWEST	72461	9.898850561769162
CENTRAL	63264	23.466578376496436
NEWTON	61160	13.979683650325562
RAMPART	55611	19.847575514305305
HOLLYWOOD	50958	27.846180453246344
OLYMPIC	48886	17.19463769893097
PACIFIC	42760	25.072161839219607
HOLLENBECK	41393	19.600667247278746
MISSION	40880	21.330305780965872
HARBOR	40637	14.160243637196546
NORTH HOLLYWOOD	39542	17.184089891996653
WILSHIRE	37712	16.080369127182834
NORTHEAST	37101	12.799215742840445
VAN NUYS	36080	19.92610084395629
WEST VALLEY	33694	15.344662228609367
TOPANGA	32340	7.019388768312766
FOOTHILL	32337	16.293601573548496
DEVONSHIRE	28673	16.71129045491336
+-----+-----+-----+		

only showing top 20 rows

Απόσταση από το πλησιέστερο αστυνομικό τμήμα:

(a)

+-----+-----+-----+		
+-----+-----+-----+		
Year	num_crimes	average_distance
+-----+-----+-----+		
2010	8213	3.965480506097993
2011	7232	2.46181888566459
2012	6550	37.04806556244542
2013	5838	2.4561803379459084
2014	4230	11.240705060052028
2015	6763	2.38790278176303
2016	8100	2.4291509215379303
2017	7788	5.620278866952371

```

|2018|    7413| 2.409083506096955|
|2019|    7129|2.4301661049761214|
|2020|    8491| 8.305664894299344|
|2021|   9767|30.666116941658924|
|2022|   10025|2.3129679282459743|
|2023|    8741|2.2716948056968675|
+-----+-----+
(b)
+-----+-----+
|    DIVISION|num_crimes| average_distance|
+-----+-----+
|    77TH STREET|    78830|1.6735955739672674|
|    SOUTHWEST|    78068| 2.161146839862122|
|    HOLLYWOOD|    70652|1.9193991303646156|
|    SOUTHEAST|    66697| 2.2223281229674|
|    OLYMPIC|   60553|1.6657781632089452|
|    CENTRAL|    59420| 0.866802366302343|
|    WILSHIRE|    58032|2.4783030061160027|
|    RAMPART|    56297| 1.361662592838379|
|    VAN NUYS|    55252|2.8073800912108116|
|    NEWTON|    45398|1.5998998286047477|
|    HOLLENBECK|    43128| 326.0868010778892|
|    PACIFIC|    40356| 3.845258334030655|
|NORTH HOLLYWOOD|    40174|2.5926392412493717|
|    HARBOR|    39433| 3.686751365873389|
|    FOOTHILL|    38327| 3.977261444227064|
|    WEST VALLEY|    34804| 2.850125414289692|
|    TOPANGA|    32648| 3.045398368787074|
|    NORTHEAST|    27283| 3.765925137599394|
|    MISSION|    27016|3.7840372209112974|
|WEST LOS ANGELES|    22376| 2.712234758681144|
+-----+-----+
only showing top 20 rows

```

## hint & explain

### Query 3

```

1 join_strategies = ["broadcast", "merge", "shuffle_hash", "shuffle_replicate_nl"]
2
3 for strategy in join_strategies:
4     print(f"Executing query with {strategy} join strategy")
5     query3_hne_results = query3_hne.query3(crime_2015, data3, revgecoding, strategy)
6     query3_hne_results.show()

```

```

1 def query3(crime_2015, data3, revgecoding, join_strategy):
2
3     crime_zip = crime_2015.join(revgecoding.hint(join_strategy), ["LAT", "LON"], "left")
4
5     best3_zip = data3.orderBy("Estimated Median Income", ascending=False).limit(3)
6     worst3_zip = data3.orderBy("Estimated Median Income", ascending=True).limit(3)
7
8     best3_zip_list = [row['Zip Code'] for row in best3_zip.collect()]
9     worst3_zip_list = [row['Zip Code'] for row in worst3_zip.collect()]
10
11     crimes = crime_zip.filter(
12         (col("ZIPcode").isin(best3_zip_list)) |
13         (col("ZIPcode").isin(worst3_zip_list))
14     )
15

```

```

16 vict_descent_count = crimes.groupBy("Vict Descent").count().orderBy("count", ascending=False)
17
18 vict_descent_count.explain()
19
20 return vict_descent_count

```

```

1 Executing query with broadcast join strategy
2 == Physical Plan ==
3 AdaptiveSparkPlan (53)
4 +- == Final Plan ==
5   TakeOrderedAndProject (31)
6   +- * HashAggregate (30)
7     +- AQEShuffleRead (29)
8       +- ShuffleQueryStage (28), Statistics(sizeInBytes=1016.0 B, rowCount=28)
9         +- Exchange (27)
10          +- * HashAggregate (26)
11            +- * Project (25)
12              +- * BroadcastHashJoin Inner BuildRight (24)
13                :- * Project (12)
14                  : +- BatchEvalPython (11)
15                    :   +- * HashAggregate (10)
16                      :     +- AQEShuffleRead (9)
17                        :       +- ShuffleQueryStage (8), Statistics(sizeInBytes=88.4 MiB, rowCount=1.96E+5)
18                          :     +- Exchange (7)
19                            :   +- * HashAggregate (6)
20                              :     +- Union (5)
21                                :   :- * Filter (2)
22                                  :     +- Scan csv (1)
23                                    +- * Filter (4)
24                                      +- Scan csv (3)
25              +- BroadcastQueryStage (23), Statistics(sizeInBytes=8.0 MiB, rowCount=682)
26                +- BroadcastExchange (22)
27                  +- * Project (21)
28                    +- * Filter (20)
29                      +- * HashAggregate (19)
30                        +- AQEShuffleRead (18)
31                          +- ShuffleQueryStage (17), Statistics(sizeInBytes=1475.8 KiB, rowCount=3.78E+4)
32                            +- Exchange (16)
33                              +- * HashAggregate (15)
34                                +- * Filter (14)
35                                  +- Scan csv (13)
36 +- == Initial Plan ==
37   TakeOrderedAndProject (52)
38   +- HashAggregate (51)
39     +- Exchange (50)
40       +- HashAggregate (49)
41         +- Project (48)
42           +- BroadcastHashJoin Inner BuildRight (47)
43             :- Project (39)
44               : +- BatchEvalPython (38)
45                 :   +- HashAggregate (37)
46                   :     +- Exchange (36)
47                     :   +- HashAggregate (35)
48                       :     +- Union (34)
49                         :   :- Filter (32)
50                           :     +- Scan csv (1)
51                             :   +- Filter (33)
52                               :     +- Scan csv (3)
53           +- BroadcastExchange (46)
54             +- Project (45)
55               +- Filter (44)
56                 +- HashAggregate (43)
57                   +- Exchange (42)
58                     +- HashAggregate (41)
59                       +- Filter (40)
60                         +- Scan csv (13)
61
62 +-----+-----+
63 | Vict Descent|count|
64 +-----+-----+
65 | Hispanic/Latin/Me...| 1053|

```

```

66 |           White| 610|
67 |           Black| 349|
68 |           Other| 272|
69 |           Unknown| 71|
70 |       Other Asian| 46|
71 |           Korean| 4|
72 |           Chinese| 1|
73 |American Indian/A...| 1|
74 |-----+-----+
75
76 Executing query with merge join strategy
77 == Physical Plan ==
78 AdaptiveSparkPlan (62)
79 +- == Final Plan ==
80   TakeOrderedAndProject (37)
81   +- * HashAggregate (36)
82     +- AQEShuffleRead (35)
83       +- ShuffleQueryStage (34), Statistics(sizeInBytes=584.0 B, rowCount=16)
84         +- Exchange (33)
85           +- * HashAggregate (32)
86             +- * Project (31)
87               +- * SortMergeJoin Inner (30)
88                 :- * Sort (16)
89                   : +- AQEShuffleRead (15)
90                     : +- ShuffleQueryStage (14), Statistics(sizeInBytes=8.7 MiB, rowCount=1.96E+5)
91                       : +- Exchange (13)
92                         +- * Project (12)
93                           +- BatchEvalPython (11)
94                             +- * HashAggregate (10)
95                               +- AQEShuffleRead (9)
96                                 +- ShuffleQueryStage (8), Statistics(sizeInBytes=88.4 MiB, rowCount=1.96E
+5)
97                                   +- Exchange (7)
98                                     +- * HashAggregate (6)
99                                       +- Union (5)
100                                         :- * Filter (2)
101                                           : +- Scan csv (1)
102                                             +- * Filter (4)
103                                               +- Scan csv (3)
104                                     +- * Sort (29)
105                                       +- AQEShuffleRead (28)
106                                         +- ShuffleQueryStage (27), Statistics(sizeInBytes=16.0 KiB, rowCount=682)
107                                           +- Exchange (26)
108                                             +- * Project (25)
109                                               +- * Filter (24)
110                                                 +- * HashAggregate (23)
111                                                   +- AQEShuffleRead (22)
112                                                     +- ShuffleQueryStage (21), Statistics(sizeInBytes=1475.8 KiB, rowCount
=3.78E+4)
113                                                       +- Exchange (20)
114                                                         +- * HashAggregate (19)
115                                                           +- * Filter (18)
116                                                             +- Scan csv (17)
117 +- == Initial Plan ==
118   TakeOrderedAndProject (61)
119   +- HashAggregate (60)
120     +- Exchange (59)
121       +- HashAggregate (58)
122         +- Project (57)
123           +- SortMergeJoin Inner (56)
124             :- Sort (47)
125               : +- Exchange (46)
126                 : +- Project (45)
127                   : +- BatchEvalPython (44)
128                     : +- HashAggregate (43)
129                       : +- Exchange (42)
130                         +- HashAggregate (41)
131                           +- Union (40)
132                             :- Filter (38)
133                               : +- Scan csv (1)
134                                 +- Filter (39)

```

```

135         :                               +- Scan csv   (3)
136     +- Sort (55)
137         +- Exchange (54)
138             +- Project (53)
139                 +- Filter (52)
140                     +- HashAggregate (51)
141                         +- Exchange (50)
142                             +- HashAggregate (49)
143                                 +- Filter (48)
144                                     +- Scan csv   (17)
145
146 +-----+-----+
147 | Vict Descent|count|
148 +-----+-----+
149 | Hispanic/Latin/M...| 1053|
150 |           White|   610|
151 |           Black|   349|
152 |           Other|   272|
153 |         Unknown|    71|
154 |       Other Asian|    46|
155 |         Korean|     4|
156 | American Indian/A...|    1|
157 |         Chinese|    1|
158 +-----+-----+
159
160 Executing query with shuffle_hash join strategy
161 == Physical Plan ==
162 AdaptiveSparkPlan (58)
163 +- == Final Plan ==
164     TakeOrderedAndProject (35)
165     +- * HashAggregate (34)
166         +- AQEShuffleRead (33)
167             +- ShuffleQueryStage (32), Statistics(sizeInBytes=584.0 B, rowCount=16)
168                 +- Exchange (31)
169                     +- * HashAggregate (30)
170                         +- * Project (29)
171                             +- * ShuffledHashJoin Inner BuildRight (28)
172                                 :- AQEShuffleRead (15)
173                                 : +- ShuffleQueryStage (14), Statistics(sizeInBytes=8.7 MiB, rowCount=1.96E+5)
174                                 :     +- Exchange (13)
175                                 :         +- * Project (12)
176                                 :             +- BatchEvalPython (11)
177                                 :                 +- * HashAggregate (10)
178                                 :                     +- AQEShuffleRead (9)
179                                 :                         +- ShuffleQueryStage (8), Statistics(sizeInBytes=88.4 MiB, rowCount=1.96E+5)
180                                 :                             +- Exchange (7)
181                                 :                                 +- * HashAggregate (6)
182                                 :                                     +- Union (5)
183                                 :                                         :- * Filter (2)
184                                 :                                             : +- Scan csv   (1)
185                                 :                                             +- * Filter (4)
186                                 :                                                 +- Scan csv   (3)
187                             +- AQEShuffleRead (27)
188                             +- ShuffleQueryStage (26), Statistics(sizeInBytes=16.0 KiB, rowCount=682)
189                                 +- Exchange (25)
190                                     +- * Project (24)
191                                         +- * Filter (23)
192                                             +- * HashAggregate (22)
193                                                 +- AQEShuffleRead (21)
194                                                     +- ShuffleQueryStage (20), Statistics(sizeInBytes=1475.8 KiB, rowCount=3.78E
195
196 +4)
197                                     +- Exchange (19)
198                                         +- * HashAggregate (18)
199                                             +- * Filter (17)
200                                                 +- Scan csv   (16)
201
202 +- == Initial Plan ==
203     TakeOrderedAndProject (57)
204     +- HashAggregate (56)
205         +- Exchange (55)
206             +- HashAggregate (54)
207                 +- Project (53)

```

```

205 +- ShuffledHashJoin Inner BuildRight (52)
206 :- Exchange (44)
207 : +- Project (43)
208 :   +- BatchEvalPython (42)
209 :     +- HashAggregate (41)
210 :       +- Exchange (40)
211 :         +- HashAggregate (39)
212 :           +- Union (38)
213 :             :- Filter (36)
214 :               : +- Scan csv (1)
215 :                 +- Filter (37)
216 :                   +- Scan csv (3)
217 +- Exchange (51)
218   +- Project (50)
219     +- Filter (49)
220       +- HashAggregate (48)
221         +- Exchange (47)
222           +- HashAggregate (46)
223             +- Filter (45)
224               +- Scan csv (16)
225
226 +-----+
227 | Vict Descent|count|
228 +-----+
229 |Hispanic/Latin/Me...| 1053|
230 |      White|      610|
231 |      Black|      349|
232 |      Other|      272|
233 |    Unknown|       71|
234 |    Other Asian|      46|
235 |      Korean|        4|
236 |     Chinese|        1|
237 |American Indian/A...|        1|
238 +-----+
239
240 Executing query with shuffle_replicate_nl join strategy
241 == Physical Plan ==
242 AdaptiveSparkPlan (50)
243 +- == Final Plan ==
244   TakeOrderedAndProject (29)
245   +- * HashAggregate (28)
246     +- AQEShuffleRead (27)
247       +- ShuffleQueryStage (26), Statistics(sizeInBytes=3.5 KiB, rowCount=100)
248         +- Exchange (25)
249           +- * HashAggregate (24)
250             +- * Project (23)
251               +- CartesianProduct Inner (22)
252                 :- * Project (12)
253                   : +- BatchEvalPython (11)
254                     :   +- * HashAggregate (10)
255                       :     +- AQEShuffleRead (9)
256                         :       +- ShuffleQueryStage (8), Statistics(sizeInBytes=88.4 MiB, rowCount=1.96E+5)
257                           :     +- Exchange (7)
258                             :   +- * HashAggregate (6)
259                               :     +- Union (5)
260                                 :- * Filter (2)
261                                   :   +- Scan csv (1)
262                                   :   +- * Filter (4)
263                                   :     +- Scan csv (3)
264             +- * Project (21)
265               +- * Filter (20)
266                 +- * HashAggregate (19)
267                   +- AQEShuffleRead (18)
268                     +- ShuffleQueryStage (17), Statistics(sizeInBytes=1475.8 KiB, rowCount=3.78E+4)
269                       +- Exchange (16)
270                         +- * HashAggregate (15)
271                           +- * Filter (14)
272                             +- Scan csv (13)
273 +- == Initial Plan ==
274   TakeOrderedAndProject (49)
275   +- HashAggregate (48)

```

```

276 +- Exchange (47)
277 +- HashAggregate (46)
278 +- Project (45)
279 +- CartesianProduct Inner (44)
280 :- Project (37)
281 : +- BatchEvalPython (36)
282 : +- HashAggregate (35)
283 : +- Exchange (34)
284 : +- HashAggregate (33)
285 : +- Union (32)
286 : :- Filter (30)
287 : +- Scan csv (1)
288 : +- Filter (34)
289 : +- Scan csv (3)
290 +- Project (43)
291 +- Filter (42)
292 +- HashAggregate (41)
293 +- Exchange (40)
294 +- HashAggregate (39)
295 +- Filter (38)
296 +- Scan csv (13)
297
298 +-----+
299 | Vict Descent|count|
300 +-----+
301 |Hispanic/Latin/Me...| 1053|
302 | White| 610|
303 | Black| 349|
304 | Other| 272|
305 | Unknown| 71|
306 | Other Asian| 46|
307 | Korean| 4|
308 | Chinese| 1|
309 |American Indian/A...| 1|
310 +-----+

```

## Query 4

```

1 def query4(df, data5, join_strategy):
2
3     def haversine(lat1, lon1, lat2, lon2):
4         # Radius of the Earth in kilometers
5         R = 6371.0
6
7         lat1_rad = math.radians(lat1)
8         lon1_rad = math.radians(lon1)
9         lat2_rad = math.radians(lat2)
10        lon2_rad = math.radians(lon2)
11
12        dlat = lat2_rad - lat1_rad
13        dlon = lon2_rad - lon1_rad
14
15        a = math.sin(dlat / 2)**2 + math.cos(lat1_rad) * math.cos(lat2_rad) * math.sin(dlon / 2)**2
16        c = 2 * math.atan2(math.sqrt(a), math.sqrt(1 - a))
17
18        distance = R * c
19        return distance
20
21    def get_distance(lat1, long1, lat2, long2):
22
23        def is_valid_coordinate(lat, lon):
24            return -90 <= lat <= 90 and -180 <= lon <= 180
25
26
27        if not is_valid_coordinate(lat1, long1) or not is_valid_coordinate(lat2, long2):
28            # Print the invalid rows
29            print(f"Invalid row: lat1={lat1}, long1={long1}, lat2={lat2}, long2={long2}")
30            return -1
31

```



```

32     try:
33         return haversine(lat1, long1, lat2, long2)
34     except ValueError:
35         return -1
36
37 df = df.filter(
38     (df["AREA NAME"] != "Null Island") &
39     (df["Weapon Used Cd"].substr(1, 1) == "1")
40 )
41
42 joined_df = df.join(data5.hint(join_strategy), df["AREA"] == data5["PREC"])
43
44 joined_df.explain()
45
46 distance_udf = udf(get_distance)
47
48 distance_df = joined_df.withColumn(
49     "DISTANCE",
50     distance_udf(
51         F.col("LAT"), F.col("LON"),
52         F.col("Y"), F.col("X")
53     ).cast("double")
54 )
55
56 query_4_1a = distance_df.groupBy("Year").agg(
57     F.avg("DISTANCE").alias("average_distance"),
58     F.count("*").alias("#")
59 ).orderBy("Year").withColumnRenamed("Year", "year")
60
61 query_4_1b = distance_df.groupBy("DIVISION").agg(
62     F.avg("DISTANCE").alias("average_distance"),
63     F.count("*").alias("#")
64 ).orderBy(F.desc("#")).withColumnRenamed("DIVISION", "division")
65
66 print("Απόσταση από το αστυνομικό τμήμα που ανέλαβε την έρευνα για το περιστατικό:")
67 print("(a)")
68 query_4_1a.show()
69 print("(b)")
70 query_4_1b.show()
71
72 cross_joined_df = df.crossJoin(data5.withColumnRenamed("LAT", "Y").withColumnRenamed("LON", "X"))
73
74 cross_joined_df = cross_joined_df.withColumn(
75     "DISTANCE",
76     distance_udf(col("LAT"), col("LON"), col("Y"), col("X")).cast("double")
77 )
78
79 windowSpec = Window.partitionBy("DR_NO").orderBy("DISTANCE")
80
81 nearest_station_df = cross_joined_df.withColumn(
82     "row_num",
83     F.row_number().over(windowSpec)
84 ).filter(col("row_num") == 1).drop("row_num")
85
86 query_4_2a = nearest_station_df.groupBy("Year").agg(
87     F.avg("DISTANCE").alias("average_distance"),
88     F.count("*").alias("#")
89 ).orderBy("Year").withColumnRenamed("Year", "year")
90
91 query_4_2b = nearest_station_df.groupBy("DIVISION").agg(
92     F.avg("DISTANCE").alias("average_distance"),
93     F.count("*").alias("#")
94 ).orderBy(F.desc("#")).withColumnRenamed("DIVISION", "division")
95
96
97 print("Απόσταση από το πλησιέστερο αστυνομικό τμήμα:")
98 print("(a)")
99 query_4_2a.show()
100 print("(b)")
101 query_4_2b.show()

```
